One Studio. Every Model. Zero Node Spaghetti.

The all-in-one local generative suite — images, video, music, voice, 3D, film & training — that just works.

■ HERO IMAGE — **hero_banner.png** — a bold 1920×1080 banner: the app window on the left, a montage of a generated portrait → animated video → 3D model → LoRA grid flowing across the right. Place at the very top.

You know the drill. You open ComfyUI. You drag in forty nodes. You wire them. One connection is wrong, so nothing runs. You fix it, tweak a scheduler, run — out of memory. You downscale, run again. Twenty minutes later you finally have *one image*.

Now you want to animate it. Different workflow. Re-learn it. Then add music — that's a *different app* entirely. Then extend the clip — load *another* ComfyUI graph. Then upscale — yet another. Every step is an export, a re-import, a new tool, a new tutorial, a new way to break.

That's not creating. That's plumbing.

Imagine instead:

You generate an image. You like it — so you **edit it in place**. You turn it into a **video** with one click (the image is *already there*, no re-import). You make a **3D model** from that same image. You extend the video to any length. You add a **generated soundtrack**. You restore an old family photo, colorize a black-and-white film, swap a character in a clip, clone a voice. And when you want a recurring character, you **train a LoRA from a single photo** — the app generates the coherent dataset for you.

All in one window. Every artifact instantly reusable by every other function. Out of the box, with best-practice settings already dialed in.

This is **LTX Studio** — a 100% local platform that gathers the generative power of the **newest, most capable open models** into a single coherent GUI, and grows every week with new features, architectures and models. No cloud. No accounts. No spaghetti. Your files never leave your machine.

■ VIDEO — **intro_flow_60s.mp4** — a 60s screen recording (no talking needed): generate image → edit → image-to-video → 3D from image → add music → extend clip. One continuous take showing artifacts flowing between tabs. This is the single most important asset.

Why it's different

The old way (nodes + separate apps)	LTX Studio
40-node graphs, manual wiring, cryptic errors	Clean GUI, sensible defaults, one button
Export → re-import between every step	Artifacts shared live across every function
A different app for images, video, music, 3D	All of it in one place

Re-learn a new workflow for each task	One consistent interface for everything
OOM crashes, manual VRAM juggling	Auto VRAM tuning for 8 GB → 128 GB
You fight the tool	You focus on the creative process

Attention to detail everywhere: a quiet checkbox that silently restructures your prompt into a clean JSON blueprint before generating. Auto-selection of the right encoder for your intent. Dropdowns that read the model header and configure themselves, color-coded by VRAM. Little things that mean you don't have to be an engineer to get great results.

Membership tiers

Free — $0. Taste the full engine. A capped number of generations per day, watermarked output. Perfect to see what's possible.

■ **Supporter — $5.** Almost every feature unlocked. Output carries a small watermark. The easiest way to support development and create daily.

Creator — $14. Full features, **no watermark**, tuned for clean, brand-safe, presentable creations.

Unlimited — $19. Full creative freedom — **no artificial content filters**. You decide what you make; the tool never says no.

■ **Studio — $25.** Everything above **plus the pro toolbox**: train your own LoRAs (image / video / music), import your own models, AI Film Studio, 3D & mocap. For creators who ship.

All tiers run 100% locally on your own GPU. Watermarking & invisible provenance signing are on the near-term roadmap.

What's inside (the short list)

Images Video Music & voice Videoclips LoRA training AI Film Studio Import models 3D models Rigging & mocap Virtual worlds Public API + agent packs

Multiple state-of-the-art models in **every** category — always your choice, never locked to one. The detailed breakdown is in the posts that follow.

■ **Pick a tier, download the app, and make something in the next 10 minutes — not after your fourth OOM crash. New features drop almost every week.**

■ IMAGE — **tiers_table.png** — a clean pricing comparison graphic (optional, reinforces the tier table).

Post 1 — Images: a full darkroom, not a single button

Generate, edit, restore, compose, and lock a face — with the best open image models, all in one tab.

■ IMAGE — **images_hero_grid.png** — a 3×2 grid: photoreal portrait, cinematic scene, anime, product shot, restored old photo, colorized B&W. Shows range.

Generate with the models that matter — your pick

- **FLUX.2**, **FLUX.1** & **FLUX.1 Krea** (photoreal aesthetic), **Z-Image**, **Ernie**, **Qwen-Image**, **SD3** — native (diffusers), no ComfyUI.
- **Auto-detection** reads the model header and configures components for you; dropdowns are **color-coded by VRAM** (green/yellow/red) so you never guess.
- Live **per-step preview** — watch the image form.

Edit anything, in place

- **Inpaint / Outpaint** natively (FLUX.1 Fill, FLUX.2 Klein) — extend a canvas or repaint a region.
- **Kontext instruction edit** — tell it what to change in words.
- **Compose 2+ photos** without cut-and-paste seams (multi-reference / stitch, optional masks).
- **Prompt Builder (JSON)** and an **invisible "structure my prompt"** checkbox that quietly turns free text into a clean blueprint before generating.

Faces & restoration — the details people notice

- **PuLID face adapter**: lock a person's identity from photos. **Upload one or several** shots — it averages them for a rock-solid likeness — with advanced controls (identity strength, when to inject, prompt-adherence).
- **Face-swap** (InsightFace) on any image.
- **Restore** damaged photos and **colorize** black-and-white with the most advanced models.
- **Upscale / super-resolution**, **segmentation** (SAM2 + GroundingDINO) for precise selections.
- **Storyboard** multiple scenes as timed keyframes.

■ VIDEO — **images_edit_demo.mp4** — 30s: generate a portrait → inpaint a hat → outpaint the scene → PuLID same face in a new setting. Show it never leaving the tab.

And the killer detail: the image you just made is **already the starting frame** for video, the input for a 3D model, or a face reference for a LoRA. No export. No re-import.

Post 2 — Video: from one image to any length, with audio

Text-to-video, image-to-video, editing, character replacement, restoration — powered by LTX-2.3, native.

■ VIDEO — **video_hero_reel.mp4** — a 45s highlight reel: t2v shot, i2v from a still, a character swap, a style change (real→anime), an upscaled clip. Fast cuts, music.

Create

- **Text-to-video** and **image-to-video** (your generated image becomes the first frame automatically).
- **Audio in the generation**: sing / lip-sync to real audio, text-to-audio, HDR/EXR output.
- **Unlimited length**: a memory-bounded autoregressive chain extends clips to any duration — with **crash-resume** and **anti-drift** face/identity re-anchoring.
- **Fast / Pro / HQ** modes; new **TeaCache** speed toggle for many-step renders.
- **Keyframe interpolation** and multi-scene **storyboard** → **one continuous video, no editing.**

Edit & direct

- **Copy the motion** of any video (ControlNet: canny / depth / pose).
- **Instruction editing** ("make it snowy winter"), **masked inpaint** where the mask **tracks the object** across frames (SAM2), **style transfer** (real↔anime), **ReTake** a section, **lip-sync a talking head**.
- **Character replacement**: swap the person in a clip while keeping the scene, motion and original audio intact — one replacement image is enough.
- **MultiTalk** lip-sync to real audio, even for long videos.

Finish & restore

- **Super-resolution** (SeedVR2, FlashVSR), **frame interpolation** for higher FPS (RIFE), **per-frame face restoration**.
- **Restore & colorize old film** with the latest models.
- Real **per-step video preview** and a credible progress bar.

■ VIDEO — **video_character_replace.mp4** — a clear before/after: original clip, then the same clip with a new character, scene & audio preserved. Strongest single video proof point.

Post 3 — Music, Voice & Videoclips

Write a song, score a video, clone a voice — then turn the song into a finished music video, automatically.

■ VIDEO — **music_videoclip_demo.mp4** — 40s: type a theme → song generates (show waveform) → auto-storyboard from the lyrics → coherent music video with scene cuts on the beat.

Music

- **Songs & instrumentals** with multiple engines: HeartMuLa, **ACE-Step** and **ACE-Step 1.5** (a 5 Hz language-model planner that reasons about structure before generating).
- **Score your video** automatically (MMAudio) — sound effects & ambience matched to the footage.
- Structured **metadata** (BPM / key / time signature / language) and an optional **AI style enrichment** that turns a few tags into a rich style caption.
- Built-in **waveform player**, **spectrogram**, and a **library** of your creations.

Voice

- **Text-to-speech & voice cloning**, native: F5-TTS, XTTS-2, Fish, VibeVoice, CosyVoice, Kokoro — one shared voice library.
- Clone from a sample, then speak any text — or use the voice as a reference for singing.

Videoclips — the standout

- **"Generate music video"**: a song becomes a **coherent video that follows the lyrics** — the LLM writes a storyboard, camera directions are **auto-generated but fully editable**, and scene cuts land on the rhythm.
- Precise **lyric-to-video timing** (forced alignment), no ghosting on scene cuts, **automatic montage**.
- Drop in your own scenes for a hybrid clip.

Post 4 — Studio: train, direct a whole film, bring your own models, automate

The pro toolbox. Where a single photo becomes a trained character, and a paragraph becomes a movie.

■ VIDEO — **studio_lora_from_one_photo.mp4** — the magic trick: start with ONE photo → app generates coherent images of that character → makes videos from them → trains an image & video LoRA. Show the pipeline end to end.

Train your own LoRAs — even from a single photo

- **Image LoRA** (FLUX.1 / FLUX.2) and **Video LoRA** (LTX-2.3), plus **music LoRAs**.
- **One-or-few-photo character LoRA**: the app first generates **coherent images** of your character, then **videos** from those, then trains image and video LoRAs from the result — an auto-built dataset with subject coverage and face focus.
- Train a **video LoRA from your own existing clips** (import, auto-caption, trim, train).
- Uncensored, purpose-built **captioner (JoyCaption)** for training-grade descriptions.

AI Film Studio — a paragraph → a finished film

- Text → **screenplay** → **assets** → **sets** → **storyboard** → **montage**, with a project database and versioning.
- **Full dubbing** pipeline (translate → synthesize → re-lip-sync), visual **drag & drop timeline** with tracks, continuity control per scene, and a **one-click Auto mode** (text → master).
- Multi-LoRA composition in the storyboard, character consistency across shots.

Bring your own models & automate

- **Import your own models** and configure paths — your fine-tunes, your checkpoints.
- **Public local API** (async) exposing every function with download links + auto-start / auto-shutdown.
- **Agent packs / skills** so assistants can drive the studio: **Hermes, OpenClaw, Claude Code, Antigravity, OpenAI Codex** and more.

■ IMAGE — **studio_filmstudio_ui.png** — a screenshot of the AI Film Studio window (screenplay + storyboard + timeline). Shows depth.

Post 5 — 3D, Rigging, Mocap & Virtual Worlds

From a flat image to a rigged, animatable 3D asset — and whole explorable scenes — fast but complete.

■ VIDEO — **3d_turntable_demo.mp4** — 30s: a generated image → 3D mesh turntable (PBR) → auto-rig skeleton → a walk animation. Then a quick fly-through of a generated 3D world.

3D models

- **Image** → **3D mesh** with PBR materials / GLB export (TRELLIS.2), viewed in a native 3D viewer.
- Use the image you **just generated** — the whole chain (generate → edit → 3D) is one flow.

Rigging & animation

- **Auto-rig**: mesh → skeleton + skinning (SkinTokens), ready to animate.
- **Motion capture** from video to drive your characters.

Virtual worlds

- **Photo** → **3D scene** (WorldGen): mesh (GLB) + splats, exportable to FBX/OBJ, explorable in a web player (orbit + WASD) and a native viewer.

Integrated end-to-end example: generate an image → edit it → image-to-video → 3D model from that same image → mocap from a video → all in one app, out of the box, with best practices baked in.

FAQ

Is it really all local?

Yes — 100% on your own GPU. No cloud calls at runtime, no accounts, your files never leave your machine. (An optional token is only ever used for a one-time download of gated model weights.)

Do I need to know ComfyUI or node graphs?

No. That's the whole point. Clean GUI, best-practice defaults, one button. If you've ever rage-quit a 40-node graph, this is for you.

What GPU do I need?

It auto-tunes for cards from **8 GB to 128 GB** (dynamic VRAM management, block-swap, quantized model options). More VRAM = faster and higher-res, but small cards are supported.

What's the difference between the tiers?

Free and Supporter ($5) include a watermark; Creator ($14) removes it; Unlimited ($19) removes content restrictions; Studio ($25) adds training, model import, AI Film Studio and 3D/mocap. All are local.

Can I train a character with just a photo?

Yes. The app generates a coherent dataset (images, then videos) from your photo(s) and trains image & video LoRAs automatically.

Can I use my own models?

Yes, on the Studio tier — import checkpoints and fine-tunes and point the app at your paths.

How often is it updated?

Frequently — expect new features, models and architectures on a roughly **weekly** cadence. This is an actively growing platform, not a frozen release.

Can other tools / agents drive it?

Yes — a local API plus agent packs (Hermes, OpenClaw, Claude Code, Antigravity, Codex) expose every function programmatically.

Reddit posts (value-first, 90/10)

Publish-ready. Lead with value, mention the tool honestly, invite discussion. Check each sub's self-promo rules and flair first.

r/StableDiffusion — [Discussion]

Title: I got tired of exporting between ComfyUI, a music app, and three upscale graphs — so everything now lives in one local GUI

Body: The thing that kills my flow isn't generation quality, it's the plumbing: generate an image here, re-import it there to animate, open another app for audio, load a different graph to extend the clip. Every step is an export/import and a new workflow to remember.
I've been building a single local app where the artifacts are shared across functions — the image you generate is already the first frame for video, the input for a 3D model, and a face reference for a LoRA. No re-import. Auto VRAM tuning so I stop OOM-ing. Curious how others here handle the multi-tool shuffle — do you batch by tool, or push through per-project? (Honest note: it's a paid app with a free tier; happy to answer anything in comments.)

r/aivideo — [Discussion]

Title: Character replacement that keeps the scene, motion AND original audio — one reference image (LTX-2.3)

Body: Swapping a character in a clip usually means the background drifts or the audio's gone. Approach that's working for me: canny structure from the source video + a single replacement first-frame + an in-context LoRA, original audio muxed back. Keeps motion and scene, changes the person. What's everyone using for identity consistency across a whole clip — LoRA, face adapter, or per-frame swap?

r/singularity or r/LocalLLaMA — [Project]

Title: Train a character LoRA from a single photo — the app builds the dataset itself (images → videos → LoRA)

Body: The dataset is the bottleneck for a consistent character. Pipeline I landed on: from one photo, generate coherent images of the same identity → make short videos from those → train image + video LoRAs on that auto-built set. Runs fully local. Would love feedback on how you all handle coverage/face-focus in auto-datasets.

■ IMAGE — **reddit_thumb.png** — a single striking before/after (character replacement or restore) for the Reddit image slot.

■ Media checklist — what to create & where

Put all files in `docs\media\`. Referenced by exact filename above. Generate with LTX Studio when the GPU is free — the app can make every one of these.

File	Type	Used in	What it shows
hero_banner.png	Image 1920×1080	Main post (top)	App + montage image→video→3D→LoRA
intro_flow_60s.mp4	Video 60s	Main post	One-take flow across tabs (MOST IMPORTANT)
tiers_table.png	Image	Main post	Pricing comparison graphic (optional)
images_hero_grid.png	Image	Post 1 (Images)	3×2 range grid
images_edit_demo.mp4	Video 30s	Post 1	Generate → inpaint → outpaint → PuLID
video_hero_reel.mp4	Video 45s	Post 2 (Video)	t2v/i2v/swap/style/upscale reel
video_character_replace.mp4	Video	Post 2	Character replacement before/after
music_videoclip_demo.mp4	Video 40s	Post 3 (Music)	Song → storyboard → music video
studio_lora_from_one_photo.mp4	Video	Post 4 (Studio)	1 photo → images → videos → LoRA
studio_filmstudio_ui.png	Image (screenshot)	Post 4	AI Film Studio window
3d_turntable_demo.mp4	Video 30s	Post 5 (3D)	Image→mesh→rig→animation + world fly-through
reddit_thumb.png	Image	Reddit	Striking before/after thumbnail

Suggested creative prompts (per house rule: always fun & varied — never cats/dogs): a neon-lit alchemist brewing storms, an astronaut DJ on a floating reef, a Victorian robot painting the northern lights, a dragon barista in a cyberpunk café. Keep the samples eye-catching and diverse.